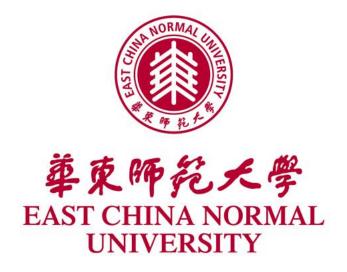
# Lock-based Concurrency Control

**SEI.ECNU** 



# **Concurrency Control**



- □ A database must provide a mechanism that will allow txns to access a database in a multi-programmed fashion while preserving the illusion that each of them is executing alone on a dedicated system.
  - serializable, recoverable and preferably cascadeless
- □ A policy in which only one transaction can execute at a time generates serial schedules, but provides a poor degree of concurrency. Why?
- Goal to develop concurrency control protocols that will assure serializability.

## **Concurrency Control Schemes**



- Two-Phase Locking (2PL)
  - Assume txns will conflict so they must acquire locks on database objects before they are allowed to access them.
- ☐ Timestamp Ordering (T/O)
  - Assume that conflicts are rare so txns do not need to first acquire locks on database objects and instead check for conflicts at commit time.

#### **Lock-Based Protocols**



- A lock is a mechanism to control concurrent access to a data item.
- Data items can be locked in two modes :
  - 1. **exclusive** (X) mode. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
  - 2. **shared** (S) mode. Data item can only be read. S-lock is requested using **lock-S** instruction.

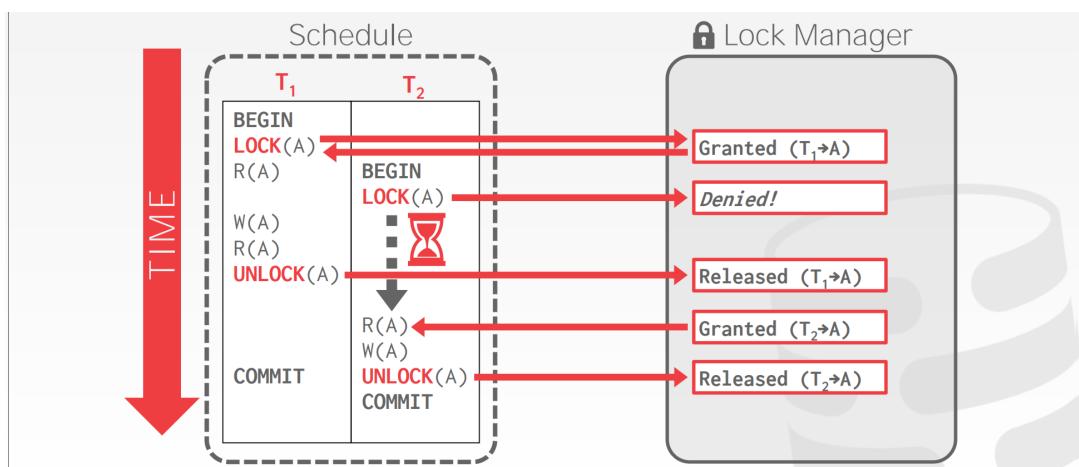
	S	X
S	true	false
X	false	false

Lock-Compatibility Matrix

- Lock requests are made to concurrency-control manager.
  - Transaction can proceed only after request is granted.

## **Executing with Locks**





#### Deadlock



- $\square$  Neither  $T_3$  nor  $T_4$  can make progress
  - Executing lock-S(B) causes T<sub>4</sub> to wait for T<sub>3</sub> to release its lock on B
  - Executing lock-X(A) causes T<sub>3</sub> to wait for T<sub>4</sub> to release its lock on A.
- ☐ Such a situation is called a **deadlock** 
  - One of  $T_3$  or  $T_4$  must be rolled back and its locks released.
- Starvation is also possible if concurrency control manager is badly designed

$T_3$	$T_4$
lock-X(B)	
read(B)	
B := B - 50	
write(B)	
	lock-S(A)
	read(A)
	lock-S(B)
lock-X(A)	

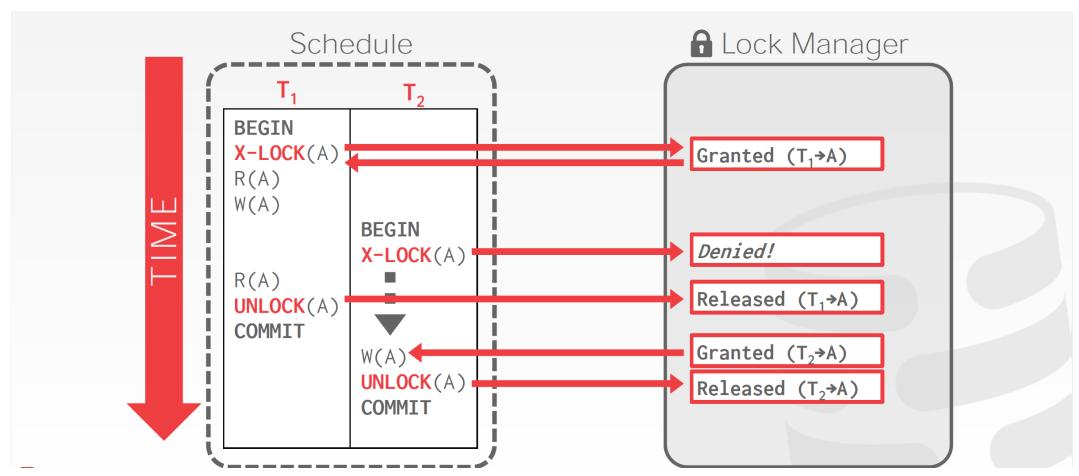
## **Two-Phase Locking**



- A protocol which ensures conflict-serializable schedules.
- □ Phase 1: Growing Phase
  - Transaction may obtain locks
  - Transaction may not release locks
- □ Phase 2: Shrinking Phase
  - Transaction may release locks
  - Transaction may not obtain locks
- □ 2PL is sufficient to guarantee conflict serializability.
  - It can be proved that the transactions can be serialized in the order of their lock points (i.e., the point where a transaction acquired its final lock).

## **Executing with 2PL**

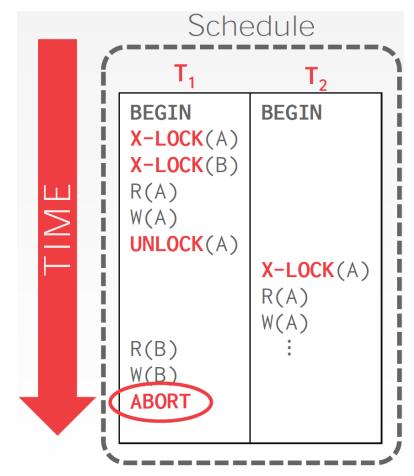




## 2PL – Cascading Aborts

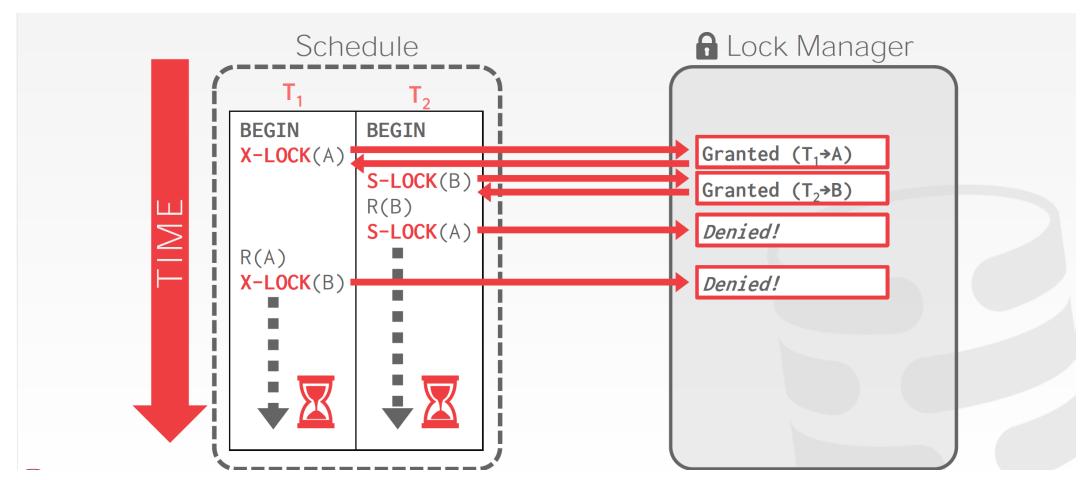


- □ This is a permissible schedule in 2PL, but the DBMS has to also abort T₂ when T₁ aborts.
  - Any information about T<sub>1</sub> cannot be "leaked" to the outside world.



#### 2PL - Deadlocks





#### **Extensions to 2PL**



- 2PL does not ensure freedom from cascading roll-back and deadlocks
- Extensions to basic 2PL
  - Strict 2PL: a transaction must hold all its exclusive locks till it commits/aborts.
    - Ensures recoverability and avoids cascading roll-backs
  - Rigorous (Strong Strict) 2PL: a transaction must hold all locks till commit/abort.
    - Transactions can be serialized in the order in which they commit.
- ☐ Most databases implement rigorous 2PL, but refer to it as 2PL

#### **Lock Conversions**



- Two-phase locking protocol with lock conversions:
  - Growing Phase:
  - can acquire a lock-S on item
  - can acquire a lock-X on item
  - can convert a lock-S to a lock-X (upgrade)
  - Shrinking Phase:
  - can release a lock-S
  - can release a lock-X
  - can convert a lock-X to a lock-S (downgrade)
- This protocol ensures serializability

#### **Automatic Acquisition of Locks**



- $\square$  A transaction  $T_i$  issues the standard read/write instruction, without explicit locking calls.
- $\square$  The operation **read**(*D*) is processed as:

```
 \begin{array}{ll} \textbf{if } T_i \text{ has a lock on } D \\ \textbf{then } & \text{read}(D) \\ \textbf{else begin} \\ & \text{if necessary wait until no other transaction has a lock-X on } D \\ & \text{grant } T_i \text{ a lock-S on } D; \\ & \text{read}(D) \\ \textbf{end} \end{array}
```

#### **Automatic Acquisition of Locks**

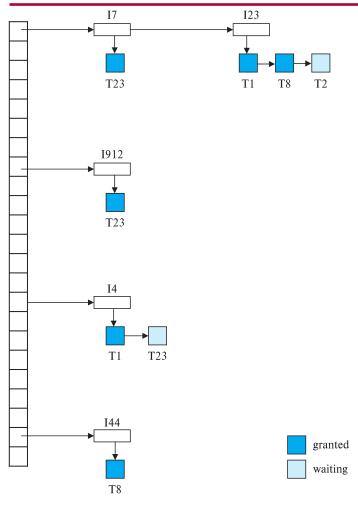


- □ All locks are released after commit or abort
- $\square$  The operation **write**(D) is processed as:

```
if T_i has a lock-X on D
        then write(D)
else begin
        if necessary wait until no other trans. has any lock on D,
        if T_i has a lock-S on D
                then upgrade lock on D to lock-X
        else grant T_i a lock-X on D
             write(D)
end;
```

#### **Lock Table**



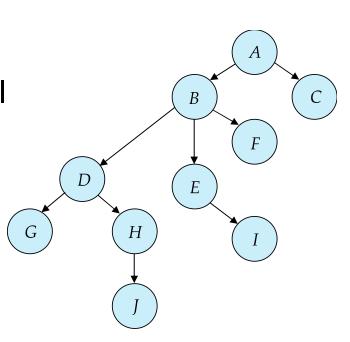


- ☐ The lock manager maintains an in-memory datastructure called a lock table to record granted locks and pending requests
  - Txns can send lock and unlock requests to LM
  - The LM replies to a lock request by sending a lock grant messages (or a message asking the transaction to roll back, in case of a deadlock)
- The LM may keep a list of locks held by each transaction

## **Graph-Based Protocols**



- An alternative to two-phase locking
- Impose a partial ordering on the dataset
- ☐ The tree-protocol is a simple kind of graph protocol
  - Ensures conflict serializability as well as freedom from deadlock
  - Unlocking may occur earlier
  - Does not guarantee recoverability or cascade freedom
    - Need to introduce commit dependencies
  - Have to lock data items that they do not access



## **Dealing with Deadlocks**



#### Approach #1: Deadlock Detection

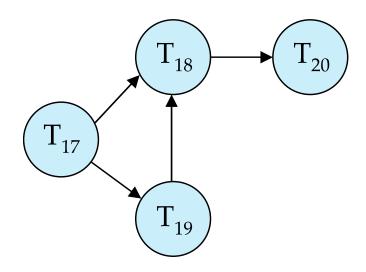
- The DBMS creates a waits-for graph to keep track of what locks each txn is waiting to acquire:
  - Nodes are transactions
  - ◆ Edge from T<sub>i</sub> to T<sub>i</sub> if T<sub>i</sub> is waiting for T<sub>i</sub> to release a lock.
- The system periodically checks for cycles in waits-for graph and then decides how to break it.

#### Approach #2: Deadlock Prevention

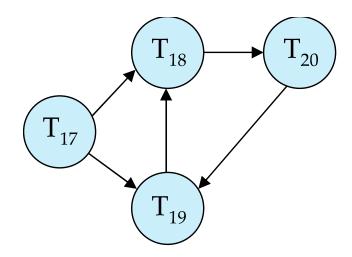
 When a txn tries to acquire a lock that is held by another txn, the DBMS kills one of them to prevent a deadlock.

#### **Deadlock Detection**





Wait-for graph without a cycle



Wait-for graph with a cycle

#### **Deadlock Prevention**



#### ■ wait-die scheme

- Older transaction may wait for younger one to release data item.
- Younger transactions never wait for older ones; they are rolled back instead.
- A transaction may die several times before acquiring a lock

#### wound-wait scheme

- Older transaction forces younger transaction rollback of instead of waiting
- Younger transactions may wait for older ones.
- Fewer rollbacks than wait-die scheme.
- In both schemes, a rolled back transactions is restarted with its original timestamp.

#### **Timeout-Based Schemes**



- ☐ A transaction waits for a lock only for a specified amount of time.
  - After that, the wait times out and the transaction is rolled back.
- Ensures that deadlocks get resolved by timeout if they occur
- ☐ Simple to implement
- But may roll back transaction unnecessarily in absence of deadlock
- Difficult to determine good value of the timeout interval.
- ☐ Starvation is also possible

## **Deadlock Recovery**



- □ Some transaction will have to rolled back (victim) to break deadlock cycle.
  - Select that transaction as victim that will incur minimum cost.
- Rollback -- determine how far to roll back transaction
  - Total rollback: Abort the transaction and then restart it.
  - Partial rollback: Roll back victim transaction only as far as necessary to release locks that another transaction in cycle is waiting for
- Starvation can happen (why?)
  - One solution: oldest transaction in the deadlock set is never chosen as victim.

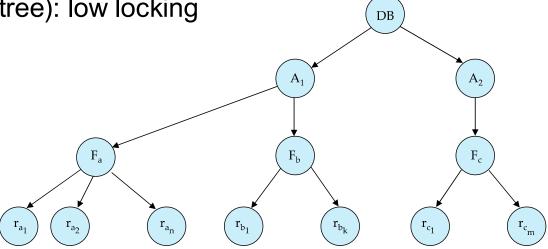
## **Multiple Granularity**



- ☐ Granularity of locking (level in tree where locking is done):
  - Fine granularity (lower in tree): high concurrency, high locking overhead

 Coarse granularity (higher in tree): low locking overhead, low concurrency

- database
- area
- file
- record



## **Lock Granularity**

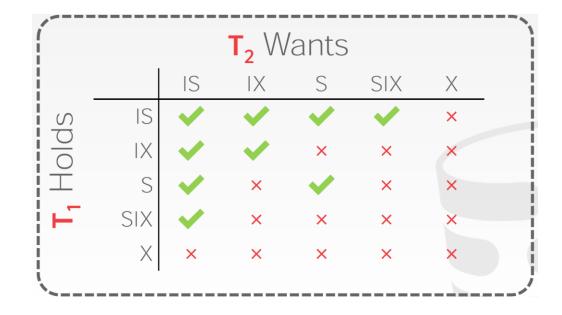


Allow data items to be of various sizes and define a hierarchy of data granularities, where the small granularities are nested within larger ones

Intention-Shared (IS): Indicates explicit locking at a lower level with shared locks.

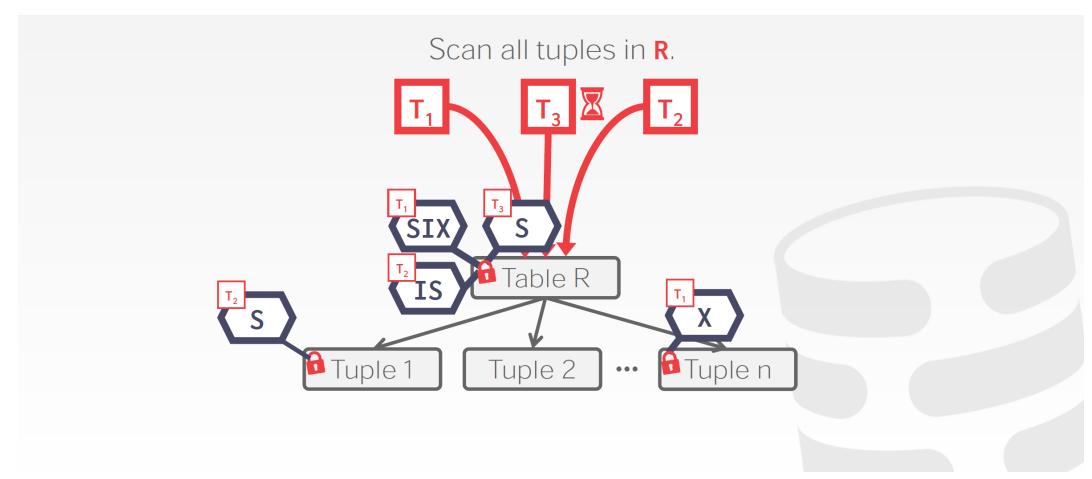
Intention-Exclusive (IX): Indicates locking at lower level with exclusive or shared locks.

Shared+Intention-Exclusive (SIX): The subtree rooted by that node is locked explicitly in shared mode and explicit locking is being done at a lower level with exclusive-mode locks.



## **Example**





## **Insert/Delete Operations**



- Locking rules for insert/delete operations
  - An X-mode lock must be obtained on an item before it is deleted
  - A transaction that inserts a new tuple into the database automatically given an X-mode lock on the tuple
- Ensures that
  - reads/writes conflict with deletes
  - Inserted tuple is not accessible by other transactions until the transaction that inserts the tuple commits

## **Predicate Reads**



T1	T2
Read(instructor where dept_name='Physics')	
	Insert Instructor in Physics
	Insert Instructor in Comp. Sci.
	Commit
Read(instructor where dept_name='Comp. Sci.')	

## **Handling Phantoms**

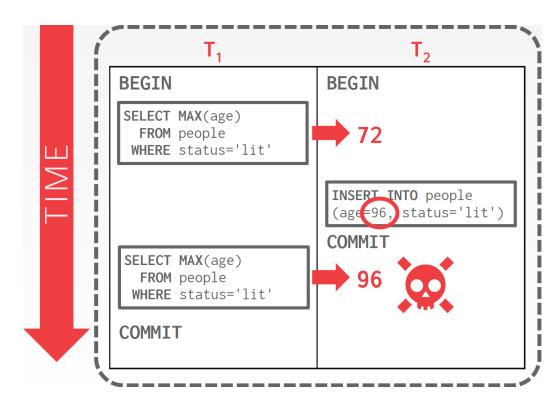


- There is a conflict at the database level
  - Predicate read or scanning the relation is reading information that indicates what tuples the relation contains
  - The transaction inserting/deleting/updating a tuple updates the same information.
  - The conflict should be detected, e.g. by locking the information.
- One solution:
  - Associate a data item with the relation, to represent the information about what tuples the relation contains.
  - Transactions scanning the relation acquire a shared lock in the data item,
  - Transactions inserting or deleting a tuple acquire an exclusive lock on the data item.
     (Note: locks on the data item do not conflict with locks on individual tuples.)
- Above protocol provides very low concurrency for insertions/deletions.

## **Predicate Locking**



- Conflict serializability on reads and writes of individual items guarantees serializability only if the set of objects is fixed.
- Lock records that satisfy a logical predicate:
  - Example: status='lit'
- ☐ In general, predicate locking has a lot of locking overhead.
- Index locking is a special case of predicate locking that is potentially more efficient.



# **Index Locking**



- Every relation must have at least one index.
- ☐ A transaction can access tuples only after finding them through one or more indices on the relation
- □ A transaction T<sub>i</sub> that performs a lookup must lock all the index leaf nodes that it accesses, in S-mode
  - Even if the leaf node does not contain any tuple satisfying the index lookup
- A transaction T<sub>i</sub> that inserts, updates or deletes a tuple t<sub>i</sub> in a relation r
  - Must update all indices to r
  - Must obtain exclusive locks on all index leaf nodes affected by the insert/update/delete
- ☐ The rules of the two-phase locking protocol must be observed

# **Next-Key Locking**



- Index-locking protocol can result in poor concurrency if there are many inserts
- Next-key locking protocol: provides higher concurrency
  - Lock all values that satisfy index lookup
  - Also lock next key value in index
    - even for inserts/deletes
  - Consider B+-tree leaf nodes as below, with query predicate 7 ≤ X ≤ 16.
     Check what happens with next-key locking when inserting: (i) 15 and (ii) 7



## **Concurrency in Index Structures**



- Indices job is to help in accessing data.
- Index-structures are typically accessed very often.
  - 2-phase locking of index nodes can lead to low concurrency.
- Index concurrency protocols
  - Locks on internal nodes are released early, and not in a two-phase fashion.
  - It is acceptable to have non-serializable concurrent access to an index
    - The accuracy of the index is maintained.
    - In particular, the exact values read in an internal node of a B<sup>+</sup>-tree are irrelevant so long
      as we land up in the correct leaf node.

## **Crabbing Protocol**



- □ Instead of two-phase locking on the nodes of the B+-tree
  - First lock the root node in shared mode.
  - After locking all required children of a node in shared mode, release the lock on the node
  - During insertion/deletion, upgrade leaf node locks to exclusive mode.
  - When splitting or coalescing requires changes to a parent, lock the parent in exclusive mode.
- Could cause excessive deadlocks
  - Can abort and restart search, without affecting transaction