

Transactions

SEI.ECNU



華東師範大學
EAST CHINA NORMAL
UNIVERSITY



Transaction Concept

- ❑ A **transaction** is a unit of program execution that accesses and possibly updates various data items.
- ❑ E.g., transaction to transfer \$50 from account A to account B:
 1. read(A)
 2. $A := A - 50$
 3. write(A)
 4. read(B)
 5. $B := B + 50$
 6. write(B)
- ❑ Two main issues to deal with:
 - **Failures** of various kinds, such as hardware failures and system crashes
 - **Concurrent** execution of multiple transactions



Requirements

□ Atomicity requirement

- If the transaction fails after step 3 and before step 6, money will be “lost” leading to an **inconsistent** database state
- Failure could be due to software or hardware
- The system should ensure that updates of a **partially** executed transaction are not reflected in the database

□ Durability requirement

- Once the user has been notified that the transaction has completed , the updates to the database by the transaction must **persist** even if there are software or hardware failures.



Requirements

□ Atomicity requirement

- If the transaction fails after step 3 and before step 6, money will be “lost” leading to an **inconsistent** database state
- Failure could be due to software or hardware
- The system should ensure that updates of a **partially** executed transaction are not reflected in the database

□ Durability requirement

- Once the user has been notified that the transaction has completed , the updates to the database by the transaction must **persist** even if there are software or hardware failures.



Requirements

□ Consistency requirement

- The sum of A and B is **unchanged** by the execution of the transaction
- Explicitly specified integrity **constraints** such as primary keys and foreign keys
- Implicit integrity constraints
- A transaction must see a **consistent** database.
 - ◆ During transaction execution the database may be **temporarily** inconsistent.
 - ◆ When the transaction completes successfully the database must be consistent
 - ◆ Erroneous transaction logic can lead to inconsistency



Requirements

□ Isolation requirement

- If T2 is allowed to access the partially updated database, it will see an inconsistent database.
- Isolation can be ensured trivially by running transactions serially

T1	T2
1. read (A)	
2. $A := A - 50$	
3. write (A)	
	read(A), read(B), print($A+B$)
4. read (B)	
5. $B := B + 50$	
6. write (B)	



ACID Properties

□ Atomicity

- Either **all** operations of the transaction are properly reflected in the database or **none** are.

□ Consistency

- Execution of a transaction in isolation preserves the **consistency** of the database.

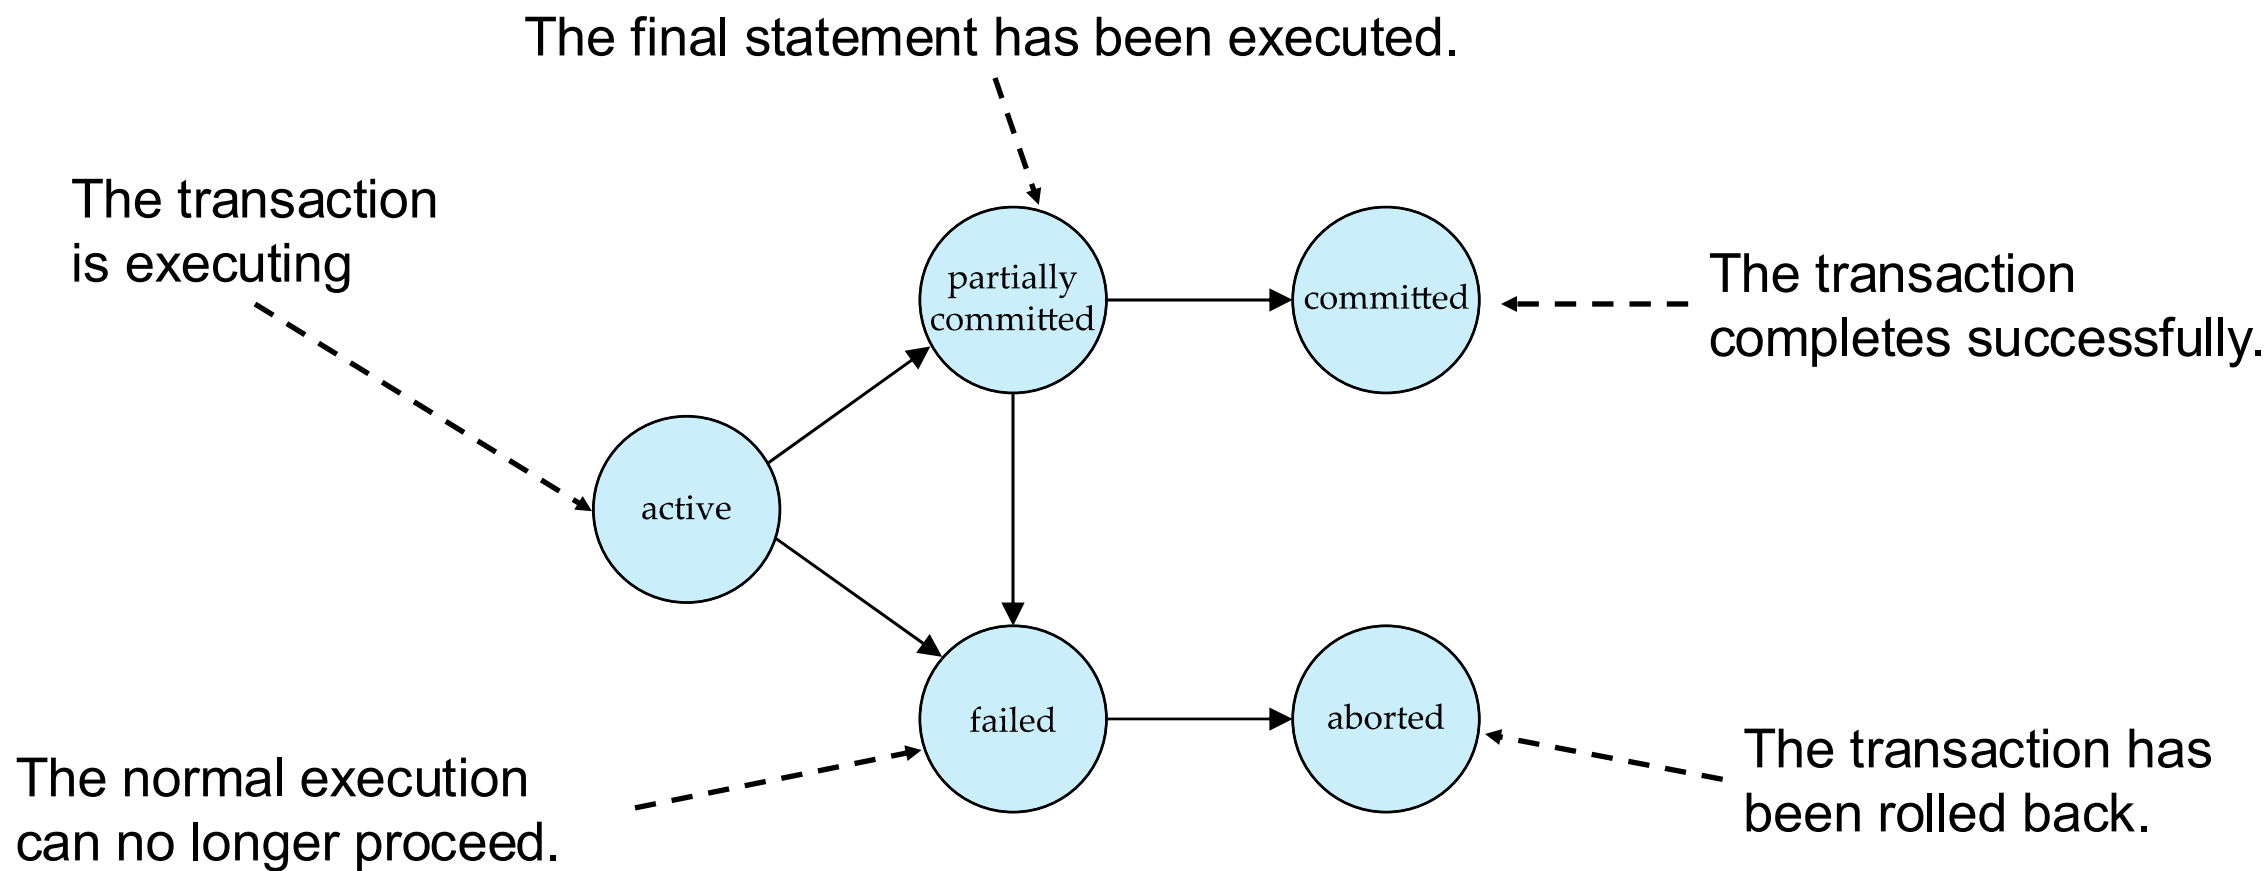
□ Isolation

- Each transaction must be **unaware** of other concurrently executing transactions.
- Intermediate transaction results must be **hidden** from other concurrently executed transactions.

□ Durability

- After a transaction completes successfully, the changes it has made to the database **persist**, even if there are system failures.

Transaction State





Concurrent Executions

□ Advantages

- **Increased processor and disk utilization**, leading to better transaction throughput
- **Reduced average response time** for transactions.

□ Concurrency control schemes

- Control the interaction among the concurrent transactions.



Schedules

- ❑ Schedule – a **sequences of instructions** that specify the chronological order in which instructions of concurrent transactions are executed
 - A schedule for a set of transactions must consist of **all instructions** of those transactions
 - Must preserve the **order** in which the instructions appear in each individual transaction.
- ❑ A transaction that successfully completes its execution will have a **commit** instructions as the last statement
 - By default transaction assumed to execute commit instruction as its last step
- ❑ A transaction that fails to successfully complete its execution will have an **abort** instruction as the last statement



Serial Schedule

T_1 : transfer \$50 from A to B ; T_2 : transfer 10% of the balance from A to B .

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

T_1	T_2
read (A) $A := A - 50$ write (A) read (B) $B := B + 50$ write (B) commit	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B) $B := B + temp$ write (B) commit

Concurrent Schedule

T_1	T_2
read (A) $A := A - 50$ write (A)	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	read (B) $B := B + temp$ write (B) commit

Equivalent to Schedule $\langle T_1, T_2 \rangle$

T_1	T_2
read (A) $A := A - 50$	read (A) $temp := A * 0.1$ $A := A - temp$ write (A) read (B)
write (A) read (B) $B := B + 50$ write (B) commit	$B := B + temp$ write (B) commit

Not equivalent to any serial schedule
Does not preserve the value of $(A + B)$.



Serializability

□ Basic Assumption

- Each transaction preserves database consistency.
- Thus, serial execution of a set of transactions preserves database consistency.

□ A (possibly concurrent) schedule is **serializable** if it is equivalent to a serial schedule

- Conflict serializability
- View serializability

- Ignore operations other than **read** and **write**
- Transactions may perform arbitrary computations on data in local buffers



Conflicting Instructions

□ Instructions I_i and I_j **conflict**

- Belong to **different** transactions
- Access the **same** data item Q
- At least one of them is **write(Q)**

1. $I_i = \text{read}(Q), I_j = \text{read}(Q).$	don't conflict.
2. $I_i = \text{read}(Q), I_j = \text{write}(Q).$	conflict.
3. $I_i = \text{write}(Q), I_j = \text{read}(Q).$	conflict
4. $I_i = \text{write}(Q), I_j = \text{write}(Q).$	conflict

□ If I_i and I_j do not conflict

- Their results would remain the same even if they had been **interchanged** in the schedule.



Conflict Serializability

□ Conflict Equivalent

- Swap non-conflicting instructions

□ Conflict Serializable

- Conflict equivalent to a serial schedule

T_3	T_4
read (Q)	
write (Q)	write (Q)

Non-Conflict Serializable

T_1	T_2
read (A)	
write (A)	read (A)
	write (A)
read (B)	
write (B)	read (B)
	write (B)

Conflict Serializable

View Serializability

□ **View Equivalent**: for each data item Q

- T_i reads the initial value of Q in both schedules
- T_i reads the result of the same $\text{Write}(Q)$ in both schedules
- The same transaction performs the final $\text{write}(Q)$ in both schedules

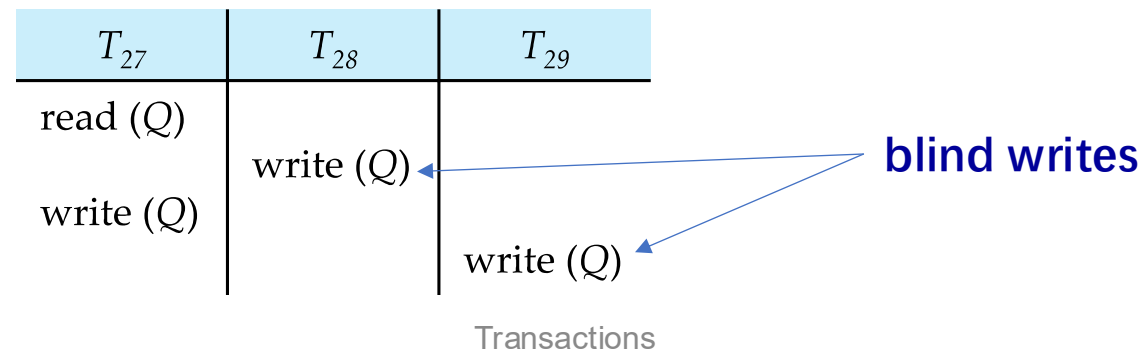
□ **View Serializable**

- View equivalent to a serial schedule
- Every conflict serializable schedule is also view serializable

T_{27}	T_{28}	T_{29}
read (Q)		
write (Q)	write (Q)	
		write (Q)

Transactions

blind writes





Other Notions of Serializability

T_1	T_5
read (A) $A := A - 50$ write (A)	
read (B) $B := B + 50$ write (B)	read (B) $B := B - 10$ write (B)
	read (A) $A := A + 10$ write (A)

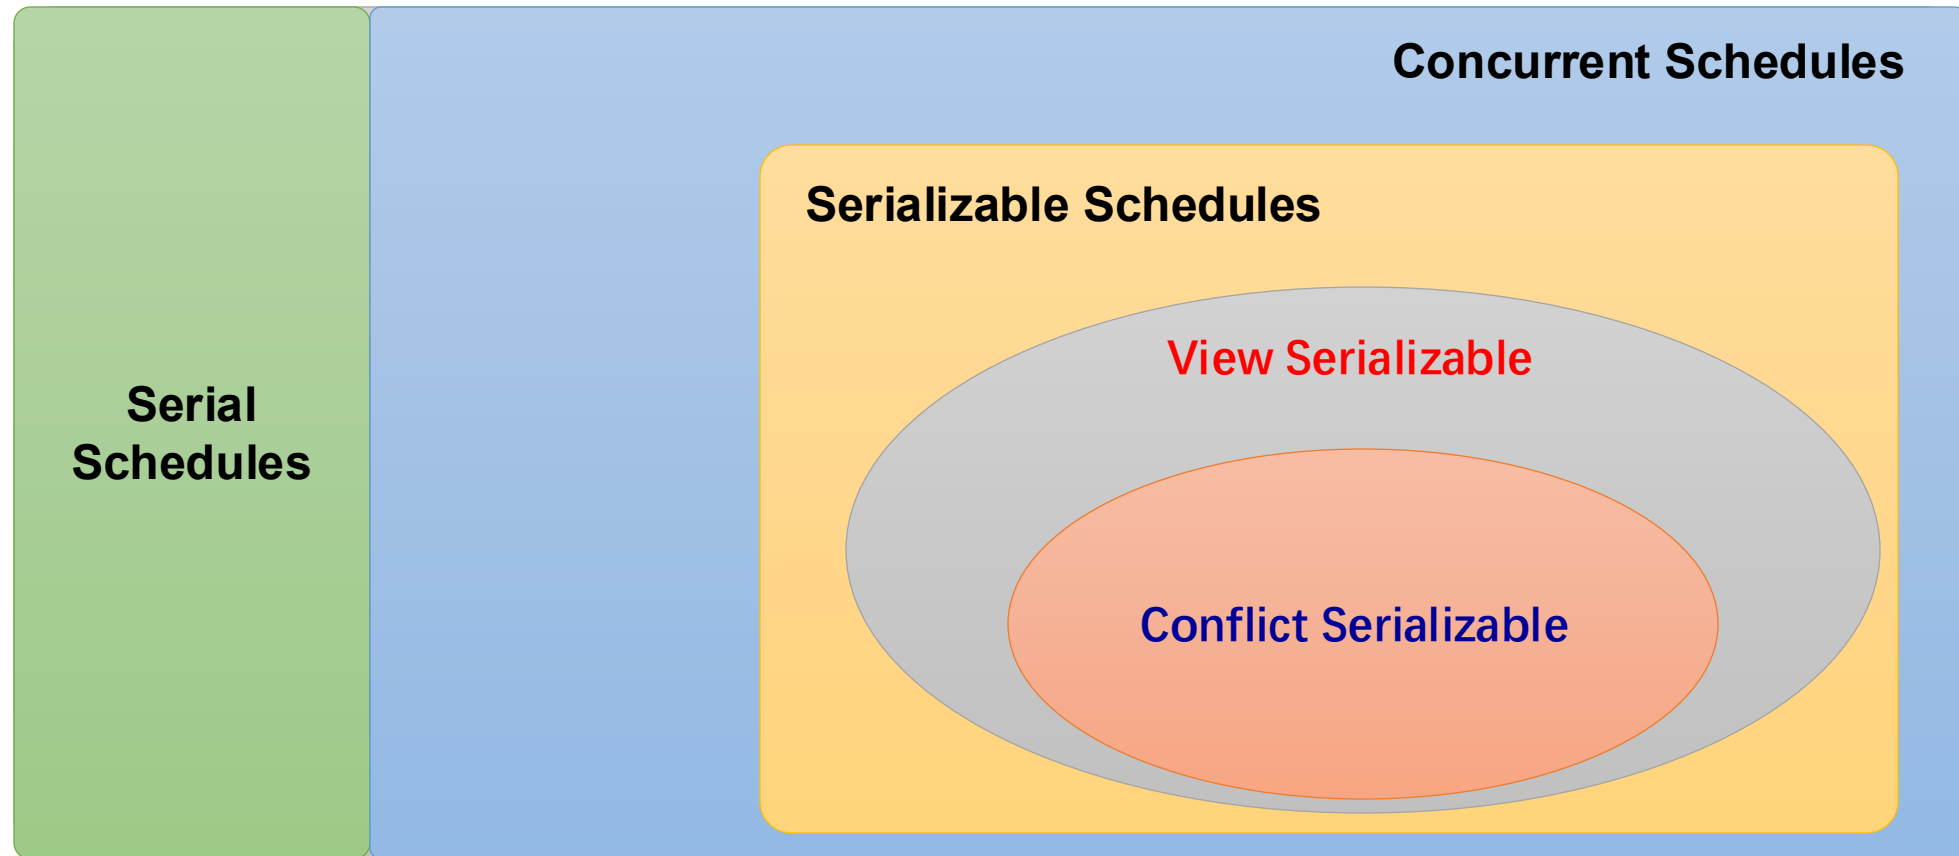
Non-Conflict Equivalent

Non-View Equivalent

Serializability




Determining such equivalence requires analysis of operations other than read and write.

Serializability



Test for Serializability

- ❑ A schedule is conflict serializable iff its precedence graph is **acyclic**.
- ❑ Test for view serializable falls in the class of **NP-complete** problems.





Recoverable Schedules

- If a transaction T_j reads a data item previously written by a transaction T_i , then the commit operation of T_i should appear before the commit operation of T_j .
- Database must ensure that schedules are recoverable

T_8	T_9
read (A)	
write (A)	
	read (A)
	commit
read (B)	

Non-Recoverable

Cascadeless Schedules

- ❑ **Cascading rollback**
- ❑ Every Cascadeless schedule is also recoverable
- ❑ It is desirable to restrict the schedules to those that are cascadeless

T_{10}	T_{11}	T_{12}
read (A) read (B) write (A)	read (A) write (A)	
abort		read (A)

If T_{10} fails, T_{11} and T_{12} must also be rolled back



Concurrency Control

- ❑ Need a mechanism to ensure all possible schedules are
 - either conflict or view **serializable**, and
 - are **recoverable** and preferably **cascadeless**
- ❑ Only one transaction can execute → Poor performance?
- ❑ Testing a schedule for serializability after it has executed
 - Too late!
 - Help us understand why a concurrency control protocol is correct
- ❑ Goal – to develop **concurrency control protocols** that will assure serializability.



Weak Levels of Consistency

- ❑ Some applications are willing to live with weak levels of consistency, allowing schedules that are **not serializable**
 - E.g., a read-only transaction that wants to get an approximate total balance of all accounts
 - E.g., database statistics computed for query optimization can be approximate (why?)
 - Such transactions need not be serializable with respect to other transactions
- ❑ **Tradeoff** accuracy for performance



Levels of Consistency in SQL-92

- ❑ **Serializable** — default
- ❑ **Repeatable read** — only committed records to be read.
 - Repeated reads of same record must return same value.
 - However, a transaction may not be serializable – it may find some records inserted by a transaction but not find others.
- ❑ **Read committed** — only committed records can be read.
 - Successive reads of record may return different (but committed) values.
- ❑ **Read uncommitted** — even uncommitted records may be read.
- ❑ **Snapshot Isolation**



Dirty Read Anomaly

- ❑ A transaction is allowed to read data from a row that has been modified by another running transaction and **not** yet committed.

Transaction 1	Transaction 2
<pre>/* Query 1 */ SELECT age FROM users WHERE id = 1; /* will read 20 */</pre>	
	<pre>/* Query 2 */ UPDATE users SET age = 21 WHERE id = 1; /* No commit here */</pre>
<pre>/* Query 1 */ SELECT age FROM users WHERE id = 1; /* will read 21 */</pre>	
	<pre>ROLLBACK; /* lock-based DIRTY READ */</pre>



Non-repeatable Read Anomaly

- During the course of a transaction, a row is retrieved twice and the values within the row **differ** between reads.

Transaction 1	Transaction 2
<pre>/* Query 1 */ SELECT age FROM users WHERE id = 1; /* will read 20 */</pre>	
	<pre>/* Query 2 */ UPDATE users SET age = 21 WHERE id = 1; COMMIT; /* in MVCC, or lock-based READ COMMITTED */</pre>
<pre>/* Query 1 */ SELECT age FROM users WHERE id = 1; /* will read 21 */</pre>	



Phantom Read Anomaly

- During the course of a transaction, **new** rows are added or removed by another transaction to the records being read.

Transaction 1	Transaction 2
<pre>/* Query 1 */ SELECT * FROM users WHERE age BETWEEN 10 AND 30;</pre>	
	<pre>/* Query 2 */ INSERT INTO users(id,name,age) VALUES (3, 'Bob', 27); COMMIT;</pre>
<pre>/* Query 1 */ SELECT * FROM users WHERE age BETWEEN 10 AND 30;</pre>	

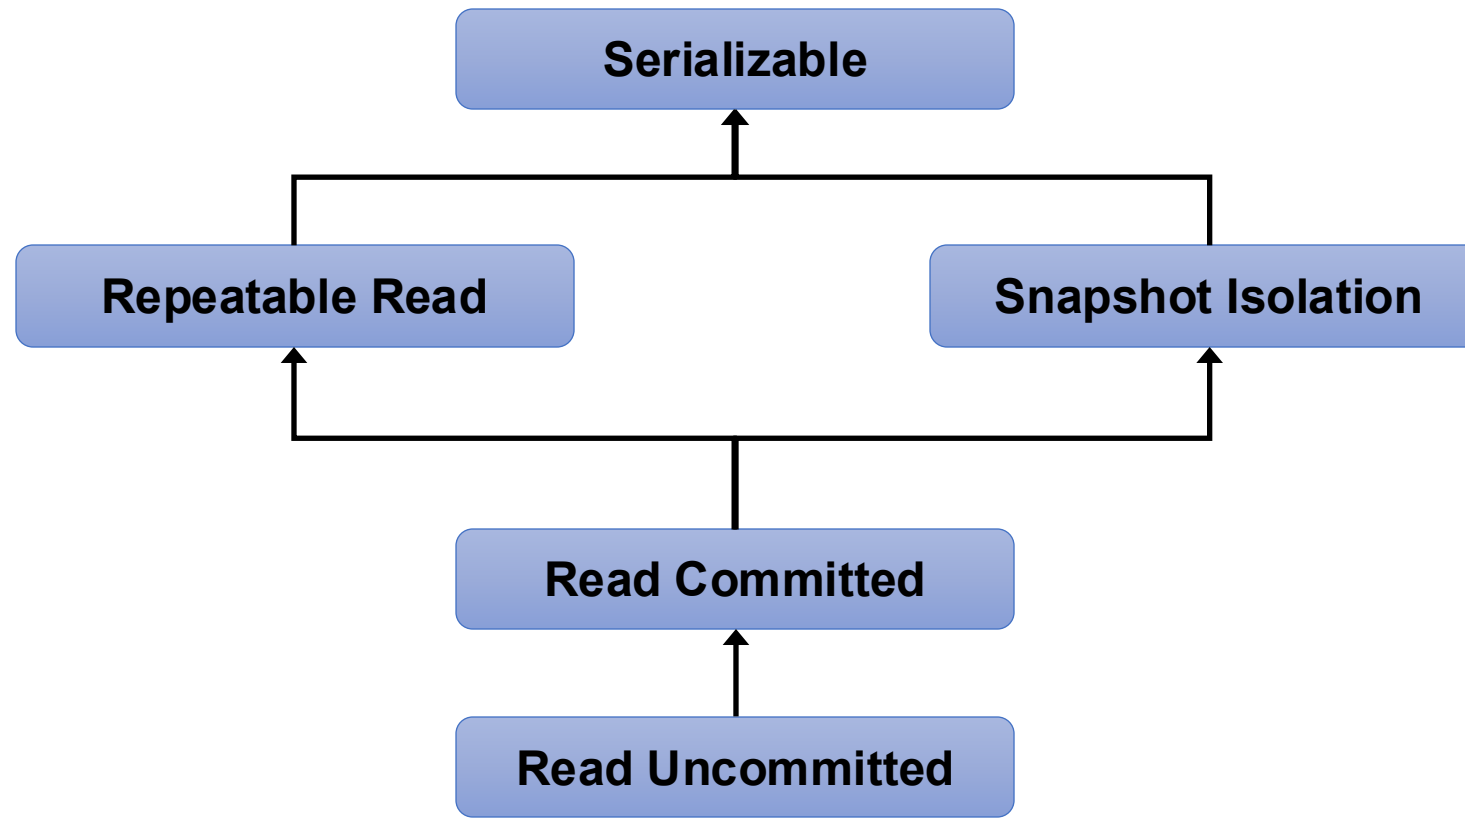


Write Skew Anomaly

- Two transactions concurrently read an **overlapping** data set, concurrently make **disjoint** updates, and finally concurrently commit, neither having seen the update performed by the other.

Transaction 1	Transaction 2
<pre>/* Keep the sum of A and B balances > 0 */ SELECT * FROM balances; /* A's balance is 100, and B's balance is 100 */</pre>	<pre>/* Keep the sum of A and B balances > 0*/ SELECT * FROM balances; /* A's balance is 100, and B's balance is 100 */</pre>
<pre>UPDATE balance SET balance = -50 WHERE id = A;</pre>	<pre>UPDATE balance SET balance = -50 WHERE id = B;</pre>
<pre>COMMIT;</pre>	<pre>COMMIT;</pre>

Isolation Level





Isolation Levels

Isolation Level	Dirty Read	Non-repeatable Read	Phantom Read	Write Skew
Read Uncommitted	✱	✱	✱	✱
Read Committed		✱	✱	✱
Repeatable Read			✱	Lock: NO MVCC: YES
Snapshot Isolation				✱
Serializable				



Transaction Definition in SQL

- ❑ In SQL, a transaction begins **implicitly**.
- ❑ A transaction in SQL ends by:
 - **Commit work / Rollback work**
- ❑ SQL statement commits implicitly by default
 - Implicit commit can be turned off by a database directive
 - ◆ JDBC -- `connection.setAutoCommit(false);`
- ❑ Isolation level can be set at (database / session / transaction) level
 - **set transaction isolation level serializable**
 - `connection.setTransactionIsolation(Connection.TRANSACTION_SERIALIZABLE)`



Implementation

□ Locking

- Lock on whole database vs lock on items
- How long to hold lock?
- Shared vs exclusive locks

□ Timestamps

- Transaction timestamp assigned e.g. when a transaction begins
- Timestamps are used to detect out of order accesses

□ Multiple versions of each data item

- Allow transactions to read from a “snapshot” of the database



Statements in OpenGauss

```
START TRANSACTION [ { ISOLATION LEVEL { READ COMMITTED | SERIALIZABLE | REPEATABLE READ }  
| { READ WRITE | READ ONLY } } [, ...] ];
```

```
BEGIN [ WORK | TRANSACTION ] [ { ISOLATION LEVEL { READ COMMITTED | SERIALIZABLE | REPEATABLE READ }  
| { READ WRITE | READ ONLY } } [, ...] ];
```

```
{ COMMIT | END } [ WORK | TRANSACTION ] ;
```

```
ROLLBACK [ WORK | TRANSACTION ];
```

```
{ SET [ LOCAL ] TRANSACTION | SET SESSION CHARACTERISTICS AS TRANSACTION }  
{ ISOLATION LEVEL { READ COMMITTED | SERIALIZABLE | REPEATABLE READ } |  
{ READ WRITE | READ ONLY } } [, ...]
```